


Cross-Cube Vega Hedging

A Swaption Market-Making Lab

Imran Hakim

Independent research — portfolio artefact
May 2026



Abstract and Executive Summary

Abstract

A swaption market-maker sees vega by cube point, but the realised cube does not move one point at a time. It moves through correlated factors: ATM level, expiry slope, curvature, and skew-wing dynamics. This lab simulates that cube, calibrates noisy SABR smiles, runs a client-flow market-maker, and compares four hedging policies. The factor-neutral rule compresses terminal P&L dispersion and improves the left tail. The limitation is explicit: no live market data, fixed beta, daily hedging, and a quoting overlay that is a factor-space heuristic rather than a full HJB solution.

Executive Summary

Single-tenor vega is a mirage. It is a convenient control number, but it is not the basis in which the cube actually moves.

The same-axis terminal P&L overlay below is the central exhibit: delta-only hedging leaves a fat distribution, while factor-neutral hedging concentrates the outcome around a tighter risk budget. The result is not free. It trades more, pays more hedge cost, and depends on liquid proxy instruments. Those costs buy a cleaner residual.

Strategy	Mean	Std	5% VaR	Sharpe
Delta-only	30.7	70.6	-77.5	0.43
Factor-neutral	9.7	6.4	-0.5	1.52




Figure 1. Same-axis terminal P&L overlay; dashed lines mark strategy means.

Introduction and Motivation

Why Node Vega Exists

Node vega is the reporting standard because it is operationally legible. Traders quote points, risk systems aggregate points, and end-of-day controls ask whether a trader is long or short a recognisable part of the cube. For a desk, that convention is useful. For a hedge, it is incomplete.

The Risk Basis

A block in 2y into 10y exposes the book to more than the 2y x 10y mark. It has projection onto level, onto the front-back expiry slope, onto an intermediate-expiry hump, and onto skew and wing dynamics. A local vega hedge can be flat at the booked point and still be long the level factor.


The Bartlett Pattern

Bartlett's SABR delta correction is a useful mental model. The Black delta is the obvious hedge only if volatility parameters stand still. Once alpha co-moves with the forward, the hedge needs the parameter response too. Cross-cube vega has the same shape: the obvious node hedge ignores nearby smiles moving together.

Trader's-Day Framing

The market-maker takes down flow, earns edge, and goes home with inventory. The next morning the cube has moved. The question is not whether the single node was hedged, but whether the hedge was short the same factor the book was long.

Architecture



Pipeline from simulated cube construction to calibration, factor analytics, market-making, and P&L attribution.

Pipeline Walk-Through

The simulator produces a full SABR cube path: forwards, annuities, strikes, ATM vols, rho, nu, alpha, and quoted smile vols. Calibration consumes a noisy snapshot and writes residual heatmaps plus static-arbitrage counts. PCA analytics recover the latent risk basis. The market-maker loop samples client flow, applies the quoting overlay, hedges at end-of-day, and emits path-level P&L. Attribution then reconciles edge, theta, delta, first-order vega, cross-vega, and hedge cost.

Reproducibility Contract

The master seed sits in YAML. The observation-noise stream is separate from the simulator streams, so calibrating a noisy market snapshot does not perturb the true cube path. The Makefile rebuilds simulations, figures, the static site, and this PDF.

Model: True Cube I

ATM Factor Surface

The cube is indexed by option expiry T in 1m, 3m, 6m, 1y, 2y, 5y, and 10y, and underlying swap tenor τ in 1y, 2y, 5y, and 10y. The ATM Black-vol surface is driven by three smooth loadings over this grid.

$$\log \sigma_{\text{ATM}}(T, \tau, t) = \mu(T, \tau) + \sum_i L_i(T, \tau) X_i(t)$$

Loading Shapes

The level loading is nearly constant across the cube with a small tenor tilt. The slope loading is the centred log-expiry coordinate. The curvature loading is a Gaussian hump around intermediate expiries, centred near two years, then normalised.

OU Dynamics

$$dX_i = -\kappa_i X_i dt + \eta_i dW_i, \quad \text{corr}(dW) = R$$

The OU factors mean-revert daily. Their instantaneous covariance is configurable in YAML, and the quoting overlay uses the full covariance matrix including off-diagonal terms. That matters because level inventory and skew inventory should interact.

Stylised Scale

The front of the normal-vol cube is set around the 60-90bp region and the back is anchored near 80bp, consistent with broad interest-rate volatility modelling examples. Absolute levels are laboratory choices, not market claims.

Model: True Cube II

Skew and Wing Dynamics

Each expiry-tenor slice carries SABR parameters alpha, beta, rho, and nu. Beta is fixed at the configured value. Rho is generated through a transformed OU process, so it stays inside $[-0.95, 0.0]$. Log-nu is another OU process. Both shocks are correlated with the level factor, reflecting the stylised observation that skew and volatility level do not move independently.

ATM Inversion

The simulator specifies ATM vol first, then recovers alpha from Hagan's ATM SABR formula. This keeps the factor model attached to the observable surface rather than to an abstract alpha process.

SABR and Vol Conversion

The pricer implements Hagan's 2002 lognormal SABR formula with a shifted extension for negative-rate laboratory cases. Black-76 and Bachelier pricing are both present. Normal/lognormal conversion is done by Brent root-finding on price equality rather than by a closed-form approximation.

Parameter Honesty

The configuration file is not a calibration file. It records a plausible set of stylised facts: mean-reverting factors, correlated level/skew shocks, bounded rho, and liquid hedge points. The report is a controlled experiment, not a claim about today's swaption market.

Model: Cube Snapshots




Figure 2. Three ATM cube snapshots from the simulated path.

Reading the Heatmaps

The snapshots are deliberately boring in the right way: the whole cube changes, and nearby expiries and tenors change together. A single isolated node shock would make the hedging problem artificial. The thesis needs a cube that moves in common modes, because that is where local vega reports become a lossy compression.

Calibration I

Noisy Market Snapshot

The calibration target is no longer the exact data-generating SABR surface. Each strike vol receives independent multiplicative lognormal observation noise: $\sigma_{\text{market}} = \sigma_{\text{true}} \exp(\epsilon)$, $\epsilon \sim N(0, \sigma_{\text{obs}}^2)$, with $\sigma_{\text{obs}} = 0.005$ by default. The random stream is separate from the simulator stream.

Per-Slice Fit

Per-slice calibration fixes beta, matches ATM exactly through alpha inversion, and fits rho and nu by least squares. It gives low local residuals, but it has no reason to respect calendar consistency across neighbouring expiries.

Cube-Constrained Fit

The joint objective keeps the same ATM discipline and adds penalties for butterfly and calendar violations, plus a smaller smoothness penalty on rho and log-nu over the cube. The no-arbitrage penalties dominate smoothness by design.

$$J = \|\sigma_m - \sigma_q\|^2 + \lambda_b B + \lambda_c C + \lambda_s S$$

This is not a local least-squares beauty contest; it is a risk-control fit. The diagnostic question is whether a small residual cost buys fewer static-arbitrage breaks.

Calibration II




Figure 3. Residual RMS in vol bp. Joint calibration pays local residual for consistency.

Static-Arbitrage Counts

Fit	Total	Butterfly	Calendar
Per-slice	3	0	3
Cube-constrained	0	0	0

The noisy per-slice fit now has genuine residuals instead of a meaningless 1e-12 colourbar. The constrained fit is slightly worse locally, but it removes the material calendar breaks under the same checker. This is the calibration analogue of the hedging thesis: a local optimum can be the wrong portfolio object.

Factor Recovery via PCA I

Method

PCA is run on daily changes in the simulated ATM log-vol surface. The first three components explain the designed level, slope, and curvature subspace. Because PCA basis vectors are sign-ambiguous, the diagnostic aligns them to the known simulator loadings and flips signs consistently.

Cosine Similarity

Factor	Cosine	Expl. var
Level	1.000	59.8%
Slope	1.000	32.8%
Curvature	1.000	7.3%

What This Validates

The table does not claim that PCA can name economic factors in real data without judgement. It says something narrower and important: in this controlled cube, the analytics recover the latent risk subspace used by the hedger.

What It Does Not Validate

It does not validate the parameter levels, the absence of jumps, or the stability of factor loadings through regimes. A live implementation would rerun this diagnostic over rolling windows.

Factor Recovery via PCA II




Figure 4. Dashed true loadings and solid recovered loadings are normalised consistently.

Reading the Figure

The recovered curves are plotted against the same expiry axis and colour as the true factor labels. Both are unit-normalised before plotting, and the sign ambiguity is resolved before the cosine table is written.

Hedging Strategies I

Delta-Only

The control hedge removes underlying swap delta and leaves vega to run. It is useful as a lower bound because it captures the cost of pretending the cube is not the main risk. Delta P&L should be small in expectation, but vega and cross-vega dominate.

Single-Tenor Vega-Flat

Each new trade's vega is snapped to the nearest cube node and offset by trading that node. This looks clean in a blotter and can be deeply wrong in factor space if the offset node loads on the wrong combination of level, slope, and curvature.

Bucketed Vega-Flat

Bucketed hedging nets vega within expiry buckets: $\leq 1y$, $1y-3y$, $3y-7y$, and $>7y$. This practical compromise cannot fully neutralise factor exposure, but it stops the most obvious front/back leakage and uses fewer hedge instruments.

Portfolio Greek Matrix

Let V be node vega and L be the loading matrix. Raw vega reports V . Factor risk is closer to $L'V$, augmented by skew exposure. The hedging question is not only whether $\text{sum}(V)=0$, but whether the projected vector is small.

Hedging Strategies II

Factor-Neutral Hedge

The factor-neutral rule selects liquid cube points and solves a small constrained hedge to bring level, slope, curvature, and skew exposures near zero. It minimises residual factor exposure subject to trading-cost and instrument-size penalties.

$$\min_h \|A(q+Bh)\|^2 + \lambda \|h\|^2 + c^\top |h|$$

In trader language: use the liquid points that actually load on the factors you are trying to offset, and avoid over-trading an ill-conditioned local hedge.

Instrument Selection

The hedge set is smaller than the risk grid. That is intentional. A market-maker normally has a handful of executable liquid points and must project the book onto that basis. The factor-neutral hedge is a practical proxy hedge, not a fantasy where every node trades frictionlessly.

Cost Accounting

Every hedge trade pays a configurable bid-ask cost in vol bp. The attribution table keeps that cost separate from cross-vega residuals.

Market-Maker Simulation

Client Flow

Client trades arrive from a Poisson process. Cube point, strike offset, side, and notional are sampled from simple distributions. The client crosses the quote with probability near one. That simplification keeps the experiment focused on inventory and hedging rather than on a separate client-response model.

Book State

After each trade, inventory Greeks and factor exposures update. Hedging fires at end-of-day. The simulation records the book before and after hedging so P&L can be attributed to edge, theta, delta, first-order vega, cross-vega, and hedge cost.

Monte Carlo Scale

The default run uses 1,000 market-making paths over 252 trading days. Tests use small deterministic seeds and avoid Monte Carlo. Generated simulations are cached under outputs/simulations, while make clean removes them to prove reproducibility.

Simplifications

There is no adverse-selection model, no funding spread, no exchange margin, and no jump risk. Those choices make the experiment narrower, but they also make the hedge comparison easier to audit.

Quoting Overlay

Reservation Vol in Factor Space

The quote mid is adjusted by an Avellaneda-Stoikov-style inventory charge, but the inventory vector is factor exposure rather than a single scalar position.

$$r_i = \text{mid}_i - \gamma q^\top \Sigma l_i (T_{\max} - t)$$




Figure 5. Half-spread responds to level and skew inventory together.

The heatmap is intentionally no longer flat. The displayed grid spans a wide enough inventory range for the risk-aversion term to matter, and the covariance matrix is used with off-diagonal terms intact. The gradient is therefore not purely axis-aligned. This remains a heuristic overlay, not a full dynamic-programming derivation.

Results I: Representative Path




Figure 6. One path shows how the same client flow diverges under different hedges.

Path-Level Read

A single path should not be over-interpreted, but it is useful for sanity. The strategies see the same simulated flow and cube path. Divergence comes from the hedge state, hedge costs, and residual factor exposures rather than from different market draws.

Results II: Distributions and Tail Risk




Figure 7. Small multiples keep the same x-axis while showing each strategy separately.

Tail-Risk Table (\$mm)

Strategy	1% VaR	5% VaR	5% CVaR	Std
Delta-only	-121.4	-77.5	-115.1	70.6
Single-tenor vega-flat	-78.9	-46.1	-62.8	37.5
Bucketed vega-flat	-55.0	-25.0	-41.3	22.3
Factor-neutral	-4.4	-0.5	-2.8	6.4

The same-axis overlay in the executive summary is the headline picture. These small multiples make the same point strategy by strategy: delta-only carries the broadest distribution, bucketed hedging improves the shape, and factor-neutral hedging is the tightest in this experiment.

Results III: P&L Attribution




Figure 8. Attribution reconciles spread edge, Greeks, residual cross-vega, and cost.

Trader-Tone Read

The factor-neutral hedge does not create edge; it protects the edge already earned. Its cost is visible and should be challenged. The reason to like it is that the cross-vega residual is less dominant, so the P&L is less dependent on being lucky about which cube factor moved after the trade.

Results IV: What the Experiment Says

Headline

The result supports the core thesis in the controlled setting: a hedge constructed in the cube's factor basis produces a tighter realised P&L distribution than a hedge constructed only in local node vega. The improvement is clearest in the left tail and in the residual attribution.

Calibration Link

The calibration section tells the same story in model space. Per-slice fitting wins the local residual contest. The cube-constrained fit is the better surface object because it removes static-arbitrage breaks under a symmetric checker.

Caveat

This is not a live trading backtest. It is a market-making lab. The aim is to make the mechanics transparent enough that a reviewer can disagree with the parameters, rerun the project, and see how the result changes.

Desk Use

The immediate desk lesson is not to throw away node vega. Keep it for operations, but report the factor projection beside it. A position that is node-flat and factor-long should not be called flat.

Assumptions and Limitations

Simulated data only. No conclusion depends on a hidden vendor feed, but absolute P&L levels are not live-market forecasts.

Single-curve discounting. The lab ignores multi-curve basis and collateral details that would matter in production pricing.

Fixed beta. SABR beta is configurable but not calibrated; this keeps the exercise focused on cross-cube vega.

No jumps. Daily OU dynamics miss event risk, central-bank gap moves, and regime shifts.

Daily hedging. Intraday inventory management is compressed into an end-of-day hedge.

Perfectly observed mids. There is no mark uncertainty or broker-quality hierarchy.

Heuristic quoting. The factor-space reservation vol borrows the Avellaneda-Stoikov shape but is not a full HJB solution.

Extensions

Real-data overlay

Use exchange settles or broker indicative marks to initialise the cube and compare factor stability.

Rough-vol dynamics

Replace OU factors with rougher latent drivers and test whether hedging cadence becomes more important.

Multi-curve pricing

Add OIS discounting and forward curves so the swaption lab better resembles production infrastructure.

Joint cap-swaption calibration

Bring caplets into the smile calibration and ask whether short-rate consistency changes the hedge basis.

Bermudan extension

Use least-squares Monte Carlo to push the factor hedge into callable-exotics inventory.

Strategic quoting

Let IDB visibility and client response feed back into the reservation-vol rule.

Learning overlay

Use reinforcement learning as a policy layer on top of the transparent factor state, not as a black-box replacement.

References I

Andersen, Leif B. G. and Vladimir V. Piterbarg (2010). *Interest Rate Modeling*. Atlantic Financial Press.

Hagan, Patrick S., Deep Kumar, Andrew S. Lesniewski, and Diana E. Woodward (2002). "Managing Smile Risk." *Wilmott Magazine*.

Bartlett, Bruce (2006). "Hedging Under SABR Model." *Wilmott Magazine*.

Rebonato, Riccardo (2002). *Modern Pricing of Interest-Rate Derivatives*. Princeton University Press.

Bergomi, Lorenzo (2016). *Stochastic Volatility Modeling*. Chapman and Hall/CRC.

References II

Avellaneda, Marco and Sasha Stoikov (2008). "High-frequency trading in a limit order book." *Quantitative Finance* 8(3), 217-224.

Gatheral, Jim (2006). *The Volatility Surface: A Practitioner's Guide*. Wiley.

Brigo, Damiano and Fabio Mercurio (2006). *Interest Rate Models: Theory and Practice*. Springer.

Glasserman, Paul (2004). *Monte Carlo Methods in Financial Engineering*. Springer.

Build Note

The preferred renderer is WeasyPrint. In this sandbox WeasyPrint's Python package installs, but native Pango/GObject libraries are unavailable, so the active renderer is headless Chromium via Playwright. Matplotlib remains available only behind the explicit emergency flag.